

Strategy, Information Architecture, and Content Architecture

Stack: Astro (SSG) · Payload CMS · Postgres · Cloudflare cache layer Goal: become India's largest packers and movers site by organic reach, without building a doorway farm.

Rev 2 — 29 July 2026. The public-branch model is removed (ADR-004 in the build plan). The site never claims physical premises: no `/branches/` pages, no branch cards, no per-city business entities. The public site speaks only of **serviceable cities and areas**. Operating bases survive as internal dispatch data that feeds the publish gate and is never rendered.

1. What we are building against

The incumbents in this vertical all fail the same way. The site we audited has 9,921 URLs in its sitemap. 7,940 of them are the same 1,086-word page with a place name swapped in, at 97% textual similarity. 95.8% of the site receives exactly one internal link. Every templated URL pattern accepts arbitrary input, so `/home-shifting-in-zzzfaketown` returns HTTP 200 with `index, follow`.

That is the competitive landscape. Nobody in Indian moving services has built a site that is both large and genuinely useful. The opportunity is not to build 50,000 pages faster than them. It is to build 6,000 pages that each contain a fact nobody else has, and to be the only site in the category that Google has no reason to distrust.

Two principles follow from that, and they govern every decision below:

1. No page ships without proprietary data. If the page can be generated from the location name alone, it does not get generated.
2. The URL space is finite and enumerable. Every valid URL comes from a Postgres row. There is no fallback route handler.

2. Positioning: trust is the product

Moving is a high-anxiety, low-frequency, high-fraud purchase. The average customer does it once every five years, is handing strangers everything they own, and has read horror stories. Consumer forums in India are full of complaints about hostage loads, quotes that double on delivery day, and companies that vanish. Roughly half the "IBA approved" claims on competitor sites are unverifiable.

So the conversion problem is not persuasion. It is risk reduction. Design and content both serve that.

What a first-time visitor must be able to verify in under 30 seconds:

SIGNAL	IMPLEMENTATION
We are a real company	Registered entity name, GSTIN, CIN in the footer on every page. Not just a phone number.
We really serve your area	Named serviceable cities and areas, each page backed by completed-job data from that area; direct phone and WhatsApp answered by a person
The price will not change	Fixed-quote guarantee stated as a policy, with the conditions that void it stated too
Someone is accountable	Named coordinator assigned per job, escalation path published, complaint SLA
Others like you used us	Reviews tied to a real job ID, a route, and a date. Not a wall of five-star quotes with first names.
Our claim process works	Published claims data: number filed, number settled, median settlement days

The last row is the one nobody else will copy, and it is the strongest trust asset we can build. Publishing an unflattering number ("median claim settlement: 11 days") is more persuasive than any testimonial.

Anti-patterns to ban at design review:

- Stock photos of smiling movers. Use photos of our crews, our trucks, our warehouses, with EXIF-plausible dates.
- Review carousels without dates, names, or verifiable job references.
- "10,000+ happy customers" style counters that can't be substantiated.
- Exit-intent popups, countdown timers, chatbots that open unprompted. All of these read as low-trust in a category already full of scams.
- Phone number as the only CTA. Offer a quote form, WhatsApp, and callback scheduling; many users won't call a stranger.
- Claiming a physical presence anywhere we do not have one. We publish no premises at all: serviceability is proven with job data, not with an address.

3. First impression and UX baseline

The homepage has one job: get the visitor to believe we are legitimate and start a quote. Everything else is secondary.

Above the fold: what we do, where we operate (detected city, confirmable), the quote entry point, and one hard credibility marker (registration + insurance + years). No hero carousel.

The quote flow is the product surface. Most competitors dump users into a generic form. Ours should be a short, progressive, three-step flow: what and where → when and size → contact. Show an indicative price range before asking for a phone number. Giving a number before taking one inverts the usual dynamic and is a trust signal in itself.

Performance targets (detail in Doc 2): LCP under 1.8s on 4G, CLS under 0.05, no layout shift from late-loading review widgets or maps. On a category where users are often on mid-range Android over patchy mobile data, speed is a trust signal.

Accessibility: WCAG 2.2 AA. Not for compliance theatre, for the practical reason that a large share of this audience is on small screens, in bright light, in a hurry, and sometimes reading in their second or third language.

4. Site structure

Keep everything within three path segments. Segment-prefixed URLs make it trivial to shard sitemaps, set cache rules, segment GSC and analytics, and apply robots policy per page family.

/	home	
/services/{service}/	national service page	~8
/packers-and-movers/{city}/	city hub	~120
/packers-and-movers/{city}/{service}/	city × service	~720
/packers-and-movers/{city}/{locality}/	locality page	~900
/routes/{origin}-to-{destination}/	lane page	~3,000
/guides/{slug}/	editorial	~150
/reviews/{city}/	city review index	~120
/tools/{slug}/	calculators, checklists	~10
/about, /pricing, /claims, /careers, ...	corporate	~25

Target at maturity: 5,000 to 7,000 pages, every one backed by data. Compare with the incumbent's 9,921 pages of which perhaps 200 are defensible.

What gets a city or locality page. Only cities we actively operate in — serviceable cities — and only localities within those cities that clear a demand threshold (search volume, or historical job count, or both). Roughly the top 15 metros × ~60 localities. Gurgaon Sector 56 gets a page. A village with 400 residents does not, no matter how tempting the long tail looks. Serviceability is decided by data (jobs completed, distance from an operating base, rate-card coverage), and the distance input stays internal — it gates publication, it is never rendered as a premises claim.

What gets a route page. Only lanes we actually operate, with a real rate card, a real transit time, and at least one completed job. Route pages are the commercial prize in this category ("Delhi to Bangalore packers and movers" is high-intent, high-value, and the incumbent doesn't even have them in its sitemap) so they deserve the most data investment, not the least.

Depth rules:

- Nothing more than 3 clicks from home.
- No page has fewer than 3 contextual inbound links from within its own topical cluster.
- No page dumps more than 80 links. City hubs paginate or group their locality lists; they do not render 664 links like the site we audited.

5. Content architecture in Payload

The core idea: pages are projections of data, not documents. An editor never writes "Packers and Movers in Rohini." They maintain the Rohini record, and the page composes itself from that record plus everything joined to it.

Collections

COLLECTION	PURPOSE	KEY FIELDS
locations	Hierarchical geo: state → city → locality	slug, type, parent, lat/lng, pincodes, population tier, is_serviceable, nearest_base
services	Service catalogue	slug, name, inclusions, exclusions, typical duration, insurance terms
lanes	Origin/destination pairs we operate	origin, destination, road_km, transit_days, rate_card, frequency, job_count
operating_bases	Internal dispatch data — never rendered. Feeds the serviceability distance in the publish gate	address, geo, licence IDs, established
rate_cards	Real pricing	service, lane or city, volume band, base, per-km, packing, insurance %, valid_from
reviews	Verified feedback	job_ref, rating, text, date, location, service, verified_by, response
jobs_stats	Aggregated operational data	location, service, month, count, on_time_%, damage_%, avg_settlement_days
faqs	Scoped Q&A	question, answer, scope (global / city / service / lane), priority
content_blocks	Reusable prose fragments	body, scope, variant_key, min_uses, max_uses
guides	Editorial	standard article model, author, reviewed_by, updated
people	Authors, ops managers, reviewers	name, role, bio, photo, credentials

Where genuine page-level variation comes from

This is the part that decides whether we are building a real site or a doorway farm. For a city × service page, the unique content is generated from data, not from spinning synonyms:

- Rate card for that service in that city, with volume bands and what's included
- Number of jobs completed there in the last 12 months, and the on-time percentage

- Localised operational facts: which societies require gate passes, parking constraints, lift availability patterns, which municipal permissions apply, typical building types
- Real reviews from that city for that service
- Route pairs most commonly booked from that city
- FAQs scoped to that city, seeded editorially and grown from actual support tickets

A locality page for Powai should mention that most towers there are 20+ floors with service lift booking rules, and that the average job is a 2BHK. A locality page for a plotted colony in Jaipur should not. That is content only an operator can write, and it is what separates us from a str_replace farm.

The editorial layer sits on top, not underneath

Every templated page has an optional `editorial_note` rich-text field. Local ops staff fill these in over time. The template renders it prominently when present. Over 18 months, the highest-traffic 500 pages become substantially hand-written, and the template becomes scaffolding rather than the whole building.

6. Programmatic SEO doctrine

The publish gate

Nothing generates unconditionally. Every page candidate is scored before it enters the build, and pages that fail are simply not built. No URL, no sitemap entry, no 404 handler that fakes one.

A reasonable scoring model (tune the weights, keep the shape):

```
score = 0
score += 25 if rate_card exists for this (service, location)
score += 20 if job_count_12m >= 10
score += 15 if reviews_count >= 3
score += 15 if operating base within 25 km (internal serviceability distance)
score += 15 if >= 200 words of location-specific editorial or ops content
score += 5 if >= 3 scoped FAQs
score += 5 if photos from this location exist
publish if score >= 60
```

Fewer than 60, the location redirects to its parent city hub. As data accumulates the page becomes eligible and appears in a later build. Growth is earned rather than declared.

Tiering

TIER	VOLUME	CONTENT MODEL	REFRESH
A — hand-written	~50	Fully editorial: home, services, pricing, claims, top 10 city hubs, cornerstone guides	Quarterly
B — template + heavy data	~1,500	City hubs, city × service, top routes. Template is 40% of the page; data and editorial are 60%	Monthly data refresh
C — template + gated data	~4,000	Localities, secondary routes. Must pass the gate	Monthly, auto-unpublish on data decay

Duplication ceiling

Enforce this in CI, not in a style guide. For every generated page, compute 8-gram Jaccard similarity against its 20 nearest siblings. Fail the build if the median exceeds 0.45. The site we audited sits at 0.68 across thousands of pages, which is what a penalty looks like before it arrives.

To keep similarity down without spinning text: rotate content_blocks by variant key, order sections by data availability rather than a fixed template, and let the amount of real data drive length. A page with a lot of local data is long. A page with less is short. Padding every page to 1,100 words is exactly the tell we're trying to avoid.

Rollout phasing

Launching 6,000 pages on day one is how a new domain gets classified as spam.

PHASE	WEEKS	SHIP	GATE TO PROCEED
0	1–6	~60 pages: home, services, 10 city hubs, guides, quote flow	Indexed, GSC clean, Core Web Vitals green
1	7–12	Top 25 cities + city × service (~200)	≥80% indexed, impressions growing, no manual action
2	13–20	Top 500 routes	≥75% indexed, routes ranking, crawl stats stable
3	21–32	Localities in metros where the gate passes	≥70% indexed
4	33+	Remaining routes and localities, 300–500 per fortnight	Rolling: indexation rate stays above 65%

The stop rule: if the indexed-to-submitted ratio for any newly shipped batch falls below 60% after four weeks, stop publishing new pages entirely and fix quality before resuming. Do not out-run indexation. Every unindexed page is Google telling us the page had no reason to exist.

7. Risks

Google's scaled content abuse policy

Since March 2024 the policy targets pages generated at scale that provide no original value, whether or not automation was involved. Intent is explicitly irrelevant. The defence is not "we used a CMS and not an LLM," it is "each page contains information that exists nowhere else and that a user would want."

Mitigations: the publish gate, the duplication ceiling, real data per page, editorial layer on high-traffic pages, phased rollout.

Doorway pages

Multiple near-identical pages targeting geographic variants, funnelling users to the same destination, is the textbook definition. We are structurally at risk because the business genuinely is one operation serving many places.

The distinguishing test: does the page deliver its own value, or is it a signpost? A locality page that lists our rate card, three reviews from that locality, its own job and on-time data, and the gate-pass rules for its main societies delivers value. One that says "we are the best packers and movers in X, call now" is a doorway. Apply this test at design review for every template.

Mitigations: every generated page must have at least one section that could not be written about any other location. If the template can't produce that section for a given location, the gate fails and the page doesn't exist.

Index bloat and crawl waste

At 6,000 pages with a bounded URL space this should not happen, but parameters, filters, and pagination will try. Handle in Doc 2 via the URL contract.

Fake reviews

India's BIS standard IS 19000:2022 covers online review integrity, and the CCPA has been increasingly active on dark patterns and fake reviews. Beyond the legal exposure, review schema abuse is a well-trodden route to losing rich results.

Rules: reviews come from real completed jobs only, tied to a job reference. Never mark up first-party reviews with AggregateRating on the organisation itself. Publish negative reviews with our response; a 4.6 average with visible 2-star reviews and replies converts better than a suspicious 4.9.

Concentration risk

A pure-SEO acquisition strategy dies with one core update. Budget for a second channel from month one: a verified Google Business Profile for the registered office (never per-city profiles for places without premises — see Doc 2 §12), a referral programme with brokers and interior firms, and the guides section built for citations rather than rankings.

Reputational risk feeding back into rankings

In this category, complaint sites rank for brand queries. Operational quality is an SEO input. A published complaints process, a fast one, protects the brand SERP.

8. Internal linking as a designed system

The incumbent's failure is instructive: 95.8% of its pages have one inbound link. Ours is computed from Postgres at build time, not hand-authored.

FROM	TO	RULE
City hub	Its localities	Grouped by zone, 8–12 per group, paginated past 60
City hub	Its city × service pages	All (6–8 links)
City hub	Top 10 routes originating there	Ranked by job volume
Locality	Parent city hub	Breadcrumb + contextual
Locality	Sibling localities	6 nearest by geo distance
Locality	Its own service pages	All
City × service	Same service in 6 nearest cities	Geo distance
Route	Both endpoint city hubs	Breadcrumb-style
Route	6 related routes	Shared origin or destination
Any page	2–3 relevant guides	Tag overlap

Every page ends up with 8 or more contextual inbound links. Anchor text comes from the data (place name, service name), varied by position, never keyword-stuffed.

Hard rule: no sitewide mega-menu listing hundreds of cities. Footer carries ~20 links maximum. Link equity flows through hierarchy, not through a nav block repeated 6,000 times.

9. Conversion architecture

Every templated page carries the same conversion spine, populated with local data:

1. Quote widget, pre-filled with the page's city or route
2. Indicative price band from the rate card, with the caveat stated honestly
3. Serviceability proof: the areas we cover in that city, jobs completed there in the last 12 months, direct phone
4. Two or three reviews scoped to that location or route
5. What's included and what costs extra, itemised
6. Scoped FAQs

7. A single closing CTA

Track quote-starts per page type, not just sessions. A locality page that gets 200 sessions and zero quote-starts is a page to rewrite or unpublish.

10. Governance

Weekly: GSC indexation by page-type segment, Core Web Vitals, crawl stats, quote-starts by template.

Monthly: duplication scores across generated families, pages with zero impressions in 90 days (candidates for unpublishing), rate card freshness, review volume per city.

Quarterly: re-run the publish gate across the whole corpus. Pages that have decayed below threshold get unpublished and 301'd to their parent. A site that prunes is a site Google trusts.

Ownership: one person owns the publish gate and the duplication ceiling and can block a release. Without that, the gate erodes within two quarters, because there is always a reason to ship one more batch of pages.

11. The one-line summary

Build 6,000 pages that each know something, ship them slowly, prune them ruthlessly, and let the competition keep their 10,000 pages of find-and-replace.